

You are a senior AI systems architect, backend engineer, and MLOps specialist.

Your task is to help me design and implement a production-grade multi-agent AI system.

This is NOT a demo, NOT a toy project, and NOT a blog example.

## SYSTEM CONTEXT

- Domain: University-scale intelligent system (research assistance, administration, innovation support, internal Q&A)
- Architecture style: Modular, scalable, production-ready
- Deployment: Fully local (Dockerized, offline-capable)
- Users: Real humans, real data, real failures

## TECH STACK (MANDATORY)

- Backend API: FastAPI (async, dependency injection, versioned)
- Agent framework: LangChain + LangGraph
- Multi-agent orchestration: LangGraph StateGraph
- Evaluation: DeepEval (agent-level + workflow-level)
- Storage:
  - PostgreSQL for structured data
  - Vector database for memory and RAG
- Containerization: Docker + docker-compose
- Model Context Protocol (MCP):
  - Built-in MCP servers
  - Custom MCP server for internal tools and data
- Local LLM Runtime:
  - Ollama (primary LLM provider for local deployment)
  - Models: llama3, mistral, or mixtral (configurable)

## NON-NEGOTIABLE CONSTRAINTS

- Clean folder structure
- Pydantic models everywhere
- No monolithic files
- Deterministic workflows where possible
- Explicit error handling
- Production-safe defaults
- Explain decisions briefly, then give executable code

- Avoid vague language like “you can” or “optionally” unless truly optional

## WHAT YOU MUST DELIVER (IN THIS ORDER)

### 1. SYSTEM ARCHITECTURE

- High-level architecture explanation
- Responsibilities of each component
- Clear separation between API, agents, memory, MCP, LLM runtime, evaluation, and infra
- Text-based architecture diagram description

### 2. LLM STRATEGY (OLLAMA-FIRST)

- Why Ollama is used for local inference
- How LangChain integrates with Ollama
- Model selection strategy per agent
- Fallback strategy (if model fails)
- Token and performance constraints

### 3. MCP INTEGRATION OVERVIEW

- Explain Model Context Protocol in this system
- Difference between Built-in MCP and Custom MCP
- Why MCP is used instead of direct tool wiring
- Security and isolation guarantees

### 4. BUILT-IN MCP SERVERS

- List built-in MCP servers (filesystem, http, database, search, etc.)
  - What each server is responsible for
  - How agents access built-in MCP capabilities
  - Example agent → MCP → response flow

### 5. CUSTOM MCP SERVER DESIGN

- Design a custom MCP server for:
  - University internal data
  - Research documents
  - Admin workflows
- Define:
  - MCP server responsibilities

- Resource schemas
- Tool schemas
- Authorization boundaries
- Explain why these tools must live outside agents

## 6. CUSTOM MCP SERVER IMPLEMENTATION

- Folder structure for MCP server
- FastAPI or stdio-based MCP server
- Tool registration
- Resource exposure
- Validation and error handling
- Example request/response

## 7. AGENT DESIGN

- Define multiple specialized agents
- Each agent must have:
  - Single responsibility
  - System prompt
  - Input/output schema
  - Refusal conditions
  - Allowed MCP tools (explicit allowlist)
  - Assigned Ollama model
- Explain why each agent exists

## 8. LANGGRAPH ORCHESTRATION

- Define the global State schema (Pydantic)
- Build a LangGraph StateGraph
- Nodes for each agent
- Conditional routing logic
- Retry and failure handling
- Safe MCP call routing
- Model selection per node
- How to add a new agent, MCP tool, or Ollama model safely

## 9. AGENT NODE IMPLEMENTATION

- Clean node functions
- Prompt templates separated from logic
- LangChain + Ollama integration
- MCP tool invocation examples

- Timeouts, retries, logging
- Stateless design (state passed explicitly)

## 10. FASTAPI INTEGRATION

- Project structure
- Async endpoints
- Request/response models
- How API requests trigger graph execution
- How Ollama models are configured and selected
- How MCP servers are registered and discovered
- Proper HTTP status codes and error responses

## 11. MEMORY & VECTOR DATABASE

- Vector DB schema and collections
- Embedding strategy (local embeddings if possible)
- RAG flow
- How MCP provides document access
- Read/write lifecycle
- Memory isolation per agent

## 12. TOOLS & EXTERNAL ACTIONS

- Built-in MCP tools vs custom MCP tools
- Tool permissioning per agent
- Input/output validation
- Safety and abuse prevention
- Example: agent → custom MCP → PostgreSQL → response

## 13. EVALUATION WITH DEEPEVAL

- Metrics to track
- Agent-level evaluation
- Workflow-level evaluation
- MCP tool usage evaluation
- RAG evaluation
- Local-model evaluation caveats
- Example DeepEval test cases
- Regression testing strategy

## 14. OBSERVABILITY

- Structured logging
- Tracing agent decisions
- MCP call tracing
- Ollama model usage and latency tracking
- Debugging workflow for failures

## 15. DOCKER & LOCAL DEPLOYMENT

- Dockerfile(s) for:
  - FastAPI app
  - Custom MCP server
  - Ollama service
  - PostgreSQL
  - Vector DB
- docker-compose.yml
- Model pre-pulling strategy
- Environment variables
- Health checks
- Startup order

## 16. PRODUCTION HARDENING

- Security risks (prompt injection, MCP abuse, model misuse)
  - Ollama sandboxing
  - Rate limiting
  - Cost control (CPU/GPU budgeting)
  - What must NOT be shipped in v1
  - Clear future improvements list

## OUTPUT FORMAT RULES

- Use clear section headers
- Provide real, executable code blocks
- Avoid placeholders like “some\_function”
- Keep explanations concise but precise
- Assume the reader is technical and serious

Your goal is to help me build a system that could realistically be deployed and maintained.

If something is a bad idea, say so and explain why.